aws

Application of Deep Learning to Text and Image Data

Module 2, Lab 5: Finetuning BERT

BERT stands for **B**idirectional **E**ncoder **R**epresentations from **T**ransformers. To learn how BERT works, let's fine-tune the **BERT** model to classify product reviews. You will use a new library called **transformers** to download a pre-trained BERT model.

You will learn:

- How to load and format the dataset
- How to load the pre-trained model
- How to train and test the model

BERT and its variants use more resources than the other models you have used so far. This may cause your instance to run out of memory. If that happens:

- Restart the kernel (Kernel->Restart from the top menu)
- Reduce the batch size
- Then re-run the code

Note: In this walkthrough, you will use a light version of the original BERT implementation called **"DistilBert"**. You can checkout the paper (https://arxiv.org/pdf/1910.01108.pdf) about it for more details.

This lab uses a dataset derived from a small sample of Amazon product reviews.

Review dataset schema:

- **reviewText:** Text of the review
- **summary:** Summary of the review

- **verified:** Whether the purchase was verified (True or False)
- **time:** UNIX timestamp for the review
- **log_votes:** Logarithm-adjusted votes log(1+votes)
- **isPositive:** Whether the review is positive or negative (1 or 0)

You will be presented with two kinds of exercises throughout the notebook: activities and challenges.

Activity

No coding is needed for an activity. You try to understand a concept,
answer questions, or run a code cell.

Challenge

Challenges are where you can practice your coding skills.

Index

In [1]:
```
%%capture

!pip install -U -q -r requirements.txt
```

Reading and formatting the dataset

First, you need to read in the product review dataset and prepare it for the BERT model. To keep the training time down, you will only use the first 2000 data points from the dataset. If you want to improve your model after you understand how to train, you can use more data to train a new model.

```
In [2]:  %%capture

         import os
         import sys
         import time
         import torch
         import pandas as pd
         from sklearn.model_selection import train_test_split
         from transformers import DistilBertForSequenceClassification, DistilBertTokeniz
         from torch.utils.data import DataLoader

         # Import system library and append path
         sys.path.insert(1, '..')

         # Setting tokenizer parallelism to false
         os.environ["TOKENIZERS_PARALLELISM"] = "false"

         # Import utility functions that provide answers to challenges
         from MLUDTI_EN_M2_Lab5_quiz_questions import *
```

Read the dataset.

```
In [3]:  df = pd.read_csv("data/NLP-REVIEW-DATA-CLASSIFICATION-TRAINING.csv")
```

Print the dataset information to see the field types.

```
In [4]:  df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 56000 entries, 0 to 55999
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  ------      --------------  -----
 0   ID          56000 non-null  int64
 1   reviewText  55989 non-null  object
 2   summary     55987 non-null  object
 3   verified    56000 non-null  bool
 4   time        56000 non-null  int64
 5   log_votes   56000 non-null  float64
 6   isPositive  56000 non-null  int64
dtypes: bool(1), float64(1), int64(3), object(2)
memory usage: 2.6+ MB
```

You do not need any of the rows that do not have **reviewText**, so drop them.

```
In [5]:  df.dropna(subset=["reviewText"], inplace=True)
```

Try it Yourself!

Activity

In [6]: `question_1`

Out[6]:

Why did you drop all rows where reviewText is empty?

This action makes the dataset the correct size for training.	If reviewText is empty for a row, then the model doesn't have anything to train on for that row.

The reviewText feature isn't used in this model, so rows with reviewText can be dropped.

| Retry | Good! You got the correct answer. |

BERT requires a lot of compute power for large datasets. To reduce the amount of time it takes to train the model, you will only use the first 2,000 data points for this lab.

In [7]:
```python
df = df.head(2000)
```

Now split the dataset into training and validation data sets, keeping 10% of the data for validation.

In [8]:
```python
# This separates 10% of the entire dataset into validation dataset.
train_texts, val_texts, train_labels, val_labels = train_test_split(
    df["reviewText"].tolist(),
    df["isPositive"].tolist(),
    test_size=0.10,
    shuffle=True,
    random_state=324,
    stratify = df["isPositive"].tolist(),
)
```

You need to tokenize the data. To do this, use a special tokenizer built for the DistilBERT model to tokenize the training and validation texts.

```
In [9]:  tokenizer = DistilBertTokenizerFast.from_pretrained('distilbert-base-uncased')

         train_encodings = tokenizer(train_texts,
                                     truncation=True,
                                     padding=True)
         val_encodings = tokenizer(val_texts,
                                   truncation=True,
                                   padding=True)
```

tokenizer_config.json: 100%   48.0/48.0 [00:00<00:00, 6.13kB/s]

vocab.txt: 100%   232k/232k [00:00<00:00, 8.36MB/s]

tokenizer.json: 100%   466k/466k [00:00<00:00, 2.99MB/s]

config.json: 100%   483/483 [00:00<00:00, 64.8kB/s]

Create a new `ReviewDataset` class to use with the BERT model. Later, you use the training and validation encoding-label pairs with this new class.

```
In [10]:  class ReviewDataset(torch.utils.data.Dataset):
              def __init__(self, encodings, labels):
                  self.encodings = encodings
                  self.labels = labels

              def __getitem__(self, idx):
                  item = {key: torch.tensor(val[idx]).to(device) for key, val in self.enc
                  item["labels"] = torch.tensor(self.labels[idx]).to(device)
                  return item

              def __len__(self):
                  return len(self.labels)

          train_dataset = ReviewDataset(train_encodings, train_labels)
          val_dataset = ReviewDataset(val_encodings, val_labels)
```

Loading the pre-trained model

Now, you need to load the model. When you do this, several warnings will print that are related to the last classification layer of BERT where you are using a randomly initialized layer. You can ignore the warnings as they are not relevant to the type of training you are doing.

```
In [11]:  model = DistilBertForSequenceClassification.from_pretrained("distilbert-base-ur
                                                                      num_labels=2)
```

model.safetensors: 100%   268M/268M [00:01<00:00, 231MB/s]

The last step is to freeze all weights until the very last classification layer in the BERT model. This helps accelerate the training process. Training the weights of the whole network (66 million weights) takes a long time. Additionally, 2000 data points would not be enough for that task. Instead, the code below freezes all the weights until the last classification layer. This means only a small portion of the weights gets updated (rest stays the same). This is a common practice with large language models.

In [12]:
```python
# Freeze the encoder weights until the classifier
for name, param in model.named_parameters():
    if "classifier" not in name:
        param.requires_grad = False
```

Training and testing the model

Now that your data is ready and you have configured your model, its time to start the fine-tuning process. This code will take **a long time** (30+ minutes) to complete with large datasets, that is why you are running it on a subset of the full review dataset.

First, define the accuracy function.

In [13]:
```python
def calculate_accuracy(output, label):
    """Calculate the accuracy of the trained network.
    output: (batch_size, num_output) float32 tensor
    label: (batch_size, ) int32 tensor """

    return (output.argmax(axis=1) == label.float()).float().mean()
```

Now you need to create the tranining and validation loop. This loop will be similar to the previous train/validation loops, however there are a few extra parameters needed due to the transformer architecture.

You need to use the `attention_mask` and get the loss from the output of the model with `loss = output[0]`

```python
In [25]: # Hyperparameters
num_epochs = 20
learning_rate = 0.005

# Get the compute device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Create data loaders
train_loader = DataLoader(train_dataset, shuffle=True,
                          batch_size=16, drop_last=True)
validation_loader = DataLoader(val_dataset, batch_size=8,
                               drop_last=True)

# Setup the optimizer
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

model = model.to(device)

for epoch in range(num_epochs):

    train_loss, val_loss, train_acc, valid_acc = 0., 0., 0., 0.

    start = time.time()
    # Training loop starts
    model.train() # put the model in training mode
    for batch in train_loader:
        # Zero the parameter gradients
        optimizer.zero_grad()
        # Put data, label and attention mask to the correct device
        data = batch["input_ids"].to(device)
        attention_mask = batch['attention_mask'].to(device)
        label = batch["labels"].to(device)

        # Make forward pass
        output = model(data, attention_mask=attention_mask, labels=label)

        # Calculate the loss (this comes from the output)
        loss = output[0]
        # Make backwards pass (calculate gradients)
        loss.backward()
        # Accumulate training accuracy and loss
        train_acc += calculate_accuracy(output.logits, label).item()
        train_loss += loss.item()
        # Update weights
        optimizer.step()

    # Validation loop:
    # This loop tests the trained network on validation dataset
    # No weight updates here
    # torch.no_grad() reduces memory usage when not training the network
    model.eval() # Activate evaluation mode
    with torch.no_grad():
        for batch in validation_loader:
            data = batch["input_ids"].to(device)
            attention_mask = batch['attention_mask'].to(device)
            label = batch["labels"].to(device)
            # Make forward pass with the trained model so far
```

```python
            output = model(data, attention_mask=attention_mask, labels=label)
            # Accumulate validation accuracy and loss
            valid_acc += calculate_accuracy(output.logits, label).item()
            val_loss += output[0].item()

    # Take averages
    train_loss /= len(train_loader)
    train_acc /= len(train_loader)
    val_loss /= len(validation_loader)
    valid_acc /= len(validation_loader)

    end = time.time()

    print("Epoch %d: train loss %.3f, train acc %.3f, val loss %.3f, val acc %.
        epoch+1, train_loss, train_acc, val_loss, valid_acc, end-start))
```

```
Epoch 1: train loss 0.448, train acc 0.810, val loss 0.414, val acc 0.825, se
conds  28.782
Epoch 2: train loss 0.435, train acc 0.810, val loss 0.411, val acc 0.795, se
conds  28.932
Epoch 3: train loss 0.428, train acc 0.817, val loss 0.390, val acc 0.835, se
conds  28.984
Epoch 4: train loss 0.419, train acc 0.818, val loss 0.378, val acc 0.850, se
conds  29.035
Epoch 5: train loss 0.411, train acc 0.824, val loss 0.375, val acc 0.885, se
conds  29.069
Epoch 6: train loss 0.407, train acc 0.825, val loss 0.390, val acc 0.810, se
conds  29.101
Epoch 7: train loss 0.396, train acc 0.824, val loss 0.356, val acc 0.860, se
conds  29.104
Epoch 8: train loss 0.395, train acc 0.820, val loss 0.350, val acc 0.870, se
conds  29.128
Epoch 9: train loss 0.386, train acc 0.833, val loss 0.342, val acc 0.880, se
conds  29.123
Epoch 10: train loss 0.381, train acc 0.841, val loss 0.339, val acc 0.875, s
econds  29.132
Epoch 11: train loss 0.379, train acc 0.830, val loss 0.332, val acc 0.880, s
econds  29.141
Epoch 12: train loss 0.377, train acc 0.833, val loss 0.329, val acc 0.870, s
econds  29.138
Epoch 13: train loss 0.376, train acc 0.831, val loss 0.323, val acc 0.885, s
econds  29.152
Epoch 14: train loss 0.372, train acc 0.842, val loss 0.320, val acc 0.875, s
econds  29.145
Epoch 15: train loss 0.360, train acc 0.845, val loss 0.319, val acc 0.870, s
econds  29.137
Epoch 16: train loss 0.360, train acc 0.843, val loss 0.322, val acc 0.890, s
econds  29.144
Epoch 17: train loss 0.357, train acc 0.842, val loss 0.314, val acc 0.880, s
econds  29.141
Epoch 18: train loss 0.360, train acc 0.843, val loss 0.311, val acc 0.890, s
econds  29.130
Epoch 19: train loss 0.358, train acc 0.845, val loss 0.311, val acc 0.890, s
econds  29.145
Epoch 20: train loss 0.352, train acc 0.842, val loss 0.306, val acc 0.880, s
econds  29.141
```

Looking at what's going on

The fine-tuned BERT model is able to correctly classify the sentiment of the most of the records in the validation set. Let's observe in more detail how the sentences are tokenized and encoded. You can do this by picking one sentence as example to look at.

```
In [15]:  st = val_texts[19]
          print(f"Sentence: {st}")
          tok = tokenizer(st, truncation=True, padding=True)
          print(f"Encoded Sentence: {tok['input_ids']}")
```

```
Sentence: An excellent resource for all scanner owners.  Seems to be up to da
te and all inclusive.  I highly recommend this product!
Encoded Sentence: [101, 2019, 6581, 7692, 2005, 2035, 26221, 5608, 1012, 384
9, 2000, 2022, 2039, 2000, 3058, 1998, 2035, 18678, 1012, 1045, 3811, 16755,
2023, 4031, 999, 102]
```

Print the vocabulary size.

```
In [16]:  # The mapped vocabulary is stored in tokenizer.vocab
          tokenizer.vocab_size
```

```
Out[16]:  30522
```

Use the encoded sentence with the tokenizer to recover the original sentence.

```
In [17]:  # Methods convert_ids_to_tokens and convert_tokens_to_ids allow to see how sent
          print(tokenizer.convert_ids_to_tokens(tok["input_ids"]))
```

```
['[CLS]', 'an', 'excellent', 'resource', 'for', 'all', 'scanner', 'owners',
'.', 'seems', 'to', 'be', 'up', 'to', 'date', 'and', 'all', 'inclusive', '.',
'i', 'highly', 'recommend', 'this', 'product', '!', '[SEP]']
```

Getting predictions on the test data

After the model is trained, you can focus on getting test data to make predictions with. Do this by:

- Reading and format the test dataset
- Passing the test data to your trained model and make predictions

In [18]:
```python
# Read the test data (It doesn't have the isPositive label)
df_test = pd.read_csv("data/NLP-REVIEW-DATA-CLASSIFICATION-TEST.csv")
df_test.head()
```

Out[18]:

	ID	reviewText	summary	verified	time	log_votes
0	33276	I've been using greeting card software for wel...	Absolutely awful.	False	1300233600	0.000000
1	20859	This version worked well for me, have upgraded...	Good for virtual machine on a mac	True	1448755200	0.000000
2	63500	Great!	Five Stars	True	1456963200	0.000000
3	4950	I can assure you that any five star review was...	SCAM	False	1400803200	2.197225
4	26509	Overall the product really seems the same but ...	Has potential but many glitches and really the...	False	1419206400	0.000000

Just as before, drop the rows that don't have the **reviewText**.

In [19]:
```python
df_test.dropna(subset=["reviewText"], inplace=True)
```

Making predictions will also take a long time with this model. To get results quickly, start by only making predictions with 15 datapoints from the test set.

In [20]:
```python
test_texts = df_test["reviewText"].tolist()[0:15]
```

In [21]:
```python
test_encodings = tokenizer(test_texts,
                           truncation=True,
                           padding=True)
```

Create labels for the test dataset to pass zeros using `[0]*len(test_texts)`.

In [22]:
```python
test_dataset = ReviewDataset(test_encodings, [0]*len(test_texts))
```

Then, create a dataloader for the test set and record the corresponding predictions.

In [23]:
```python
test_loader = DataLoader(test_dataset, batch_size=4)
test_predictions = []
model.eval()

with torch.no_grad():
    for batch in test_loader:
        data = batch["input_ids"].to(device)
        attention_mask = batch['attention_mask'].to(device)
        label = batch["labels"].to(device)
        output = model(data, attention_mask=attention_mask, labels=label)
        predictions = torch.argmax(output.logits, dim=-1)
        test_predictions.extend(predictions.cpu().numpy())
```

Finally, pick an example sentence and examine the prediction. Does the prediction look correct?

Remember

- 1->positive class
- 0->negative class

In [24]:
```python
k = 13
print(f'Text: {test_texts[k]}')
print(f'Prediction: {test_predictions[k]}')
```

```
Text: I have been using this product for five or six years.  This purchase wa
s my annual subscription renewal.  It has the features that I need, and seems
to protect the three PC's that we have while using the internet.  I will be l
ooking at the mobile apps for a tablet and smart phone.
Prediction: 1
```

Try it Yourself!

Challenge

You trained the model for 10 epochs. Would you get better results from the validation dataset if the model trained longer?

Make a note of your last `Val_loss` result.
Then, in the Training and testing the model section, change the `num_epochs` parameter to `20`.
Finally, re-run the code blocks to load the pre-trained model, and train your model.
Did `Val_loss` improve?